

Java Singleton pattern

Contents

1. [What is the Singleton pattern, why is it used, and what does it solve?](#)
2. [What are the characteristics of a Singleton?](#)
3. [What are the advantages and disadvantages?](#)
4. [When should you use it, when avoid it, and can it violate OO principles?](#)
5. [How do you implement a Singleton, and why are the members private/static?](#)
6. [Can a Singleton be inherited, final, abstract, generic, or implement interfaces?](#)
7. [What are the types of Singleton?](#)
8. [Why isn't a simple lazy Singleton thread-safe, and how do you end up with two instances?](#)
9. [How does the synchronized Singleton work, and why is it slower?](#)
10. [What is double-checked locking, why is volatile required, and what was the pre-Java-5 problem?](#)
11. [Bill Pugh v/s double-checked locking, and which has the best balance?](#)
12. [Is the Enum Singleton thread-safe, and why is it often recommended?](#)
13. [How can a Singleton's single-instance guarantee be broken, and how do you protect it?](#)
14. [How do you test components that depend on a Singleton?](#)
15. [Would you use a Singleton for a database connection? Why or why not?](#)
16. [Have you used a Singleton in production, which implementation, and have you replaced one with DI?](#)
17. [If Spring already manages singleton beans, why does the Singleton pattern still exist?](#)
18. [How would you decide between a Singleton, dependency injection, and a static utility class?](#)
19. [Is Singleton an anti-pattern, and would you keep it if redesigning Java today?](#)
20. [Common mistakes.](#)

What the pattern is, how to implement it safely, the thread-safety and reflection traps, and when to reach for dependency injection instead. The tone is how I would talk each one through with a teammate; the design and experience questions are fully spoken.

1. What is the Singleton pattern, why is it used, and what does it solve?

So the Singleton pattern makes sure a class has exactly one instance and gives everyone one shared point of access to it. You use it when having two of something would be wrong or wasteful, a single configuration holder, a single logging component, a single connection pool. The problem it solves is “there should be one and only one of this, and everyone should share it”, instead of each caller creating its own copy. And yes, it is a creational design pattern (it is about controlling how an object gets created).

2. What are the characteristics of a Singleton?

Three things define it: a private constructor (so nobody else can new it), a static reference that holds the single instance, and a public static accessor (usually `getInstance()`) that hands out that one instance. Put together, the class controls its own creation and guarantees there is only ever one.

3. What are the advantages and disadvantages?

Advantages: one shared instance (saves memory for expensive objects), a single access point, and controlled creation. Disadvantages, and these are the reason it gets a bad name: it is global state, which makes code hard to test (you cannot easily swap it for a fake), it hides dependencies (a class secretly reaches for the singleton instead of declaring what it needs), and a mutable singleton shared across threads needs careful synchronisation. So it is convenient but it fights testability and clean design.

4. When should you use it, when avoid it, and can it violate OO principles?

Use it for something genuinely single and mostly stateless: configuration, a logger, a connection pool, a cache manager. Avoid it when you are really just reaching for a global variable, or when the object has behaviour you will want to swap or mock, that is a job for dependency injection. And yes, it can violate OO principles: it often breaks the Single Responsibility Principle (the class manages both its job and its own lifecycle) and it works against dependency inversion by hard-wiring callers to a concrete global. So use it sparingly and honestly.

5. How do you implement a Singleton, and why are the members private/static?

The essential shape is three parts: a private constructor, a private static field holding the instance, and a public static `getInstance()` that returns it. On the “why” of each:

- The constructor is private so no other code can create an instance with `new`, that is what enforces the “only one”.
- The instance field is private static because static means it belongs to the class (one per class, not per object), and private stops anyone bypassing the accessor.
- The accessor is public static so callers can get the instance without already having one (you cannot call an instance method if you cannot create an instance).

6. Can a Singleton be inherited, final, abstract, generic, or implement interfaces?

Rapid-fire, because these come up as checks:

Question	Answer
Can it be inherited?	Not really, a private constructor blocks subclassing
Should it be <code>final</code> ?	Yes, a good idea, it makes the no-subclass intent explicit
Can it be abstract?	No point, an abstract class cannot be instantiated, which contradicts “create one instance”
Can it implement interfaces?	Yes, and this is actually good (lets you inject it behind an interface)
Can it be generic?	Not cleanly, a static field cannot use the class’s type parameter, so a generic singleton does not really work

7. What are the types of Singleton?

There are a handful of implementations, and the differences are all about thread safety and laziness:

Type	Thread-safe	Lazy	Notes
Eager	Yes	No	Instance created at class load; simple, but built even if never used
Simple lazy	No	Yes	Created on first call, but the null-check races across threads
Synchronized method	Yes	Yes	Correct but locks on every call, so slower
Double-checked locking	Yes	Yes	Locks only on first creation; needs volatile
Bill Pugh (holder idiom)	Yes	Yes	Static inner holder class; thread-safe and lazy with no locking
Enum	Yes	No	Simplest safe option; also immune to reflection and serialization attacks

8. Why isn’t a simple lazy Singleton thread-safe, and how do you end up with two instances?


Here is the race. The lazy version does `if (instance == null) instance = new Singleton();`. Now picture two threads calling `getInstance()` at the same time: both check `instance == null`, both see true (neither has created it yet), so both go on to create one, and you end up with two instances. The check-then-create is not atomic, so without synchronisation the “only one” guarantee is broken the moment two threads hit it together.


9. How does the synchronized Singleton work, and why is it slower?

You fix the race by making `getInstance()` a synchronized method, so only one thread can be inside it at a time, and the null-check-and-create becomes atomic. That is correct, but it is slower because every single call takes the lock, even long after the instance exists, when there is nothing to protect anymore. So you pay the locking cost forever for a race that only mattered on the very first call. That waste is exactly what double-checked locking fixes.

10. What is double-checked locking, why is volatile required, and what was the pre-Java-5 problem?

Double-checked locking (DCL) says: check without a lock first (the fast path), and only if the instance is null do you take the lock and check again before creating. So the lock is used only on the first creation, and every call after that skips it.





The `volatile` on the instance field is not optional. Without it, the JVM is allowed to reorder the steps of `instance = new Singleton()` (allocate memory, assign the reference, run the constructor), so another thread could see a non-null but half-constructed object. `volatile` forbids that reordering and guarantees visibility, so the reference is only visible once the object is fully built. And before Java 5, DCL was actually broken: the old memory model did not give `volatile` strong enough guarantees, so DCL did not work reliably until Java 5 revised the memory model.

11. Bill Pugh v/s double-checked locking, and which has the best balance?

The Bill Pugh approach (the initialization-on-demand holder idiom) puts the instance in a static inner class: `private static class Holder { static final Singleton INSTANCE = new Singleton(); }`, and `getInstance()` returns `Holder.INSTANCE`. The trick is that the JVM does not load the inner class until it is first referenced, so the instance is created lazily, and class loading is thread-safe by the JVM's own guarantee, so you get laziness and thread safety with no synchronisation and no `volatile`. Compared to DCL it is simpler (no tricky volatile, no double check to get wrong) and just as fast. So for a plain class-based singleton, Bill Pugh has the best balance of simplicity and performance.

12. Is the Enum Singleton thread-safe, and why is it often recommended?

Yes, an enum singleton is thread-safe by construction, the JVM guarantees enum values are created once, safely. It is often the recommended way (Joshua Bloch's advice) because it also handles the two things that break other singletons for free: it is immune to reflection attacks and it serializes correctly without you writing `readResolve()`. The trade-off is that it is eager (not lazy) and cannot extend another class, but for most singletons `enum Singleton { INSTANCE; ... }` is the simplest bulletproof option.

13. How can a Singleton's single-instance guarantee be broken, and how do you protect it?

Three sneaky ways beyond the threading race:

- Reflection: someone calls `setAccessible(true)` on the private constructor and creates a second instance. Defense: throw from the constructor if the instance already exists.

- Serialization: deserializing a singleton creates a new object. Defense: implement `readResolve()` to return the existing instance.
- Cloning: if the class is `Cloneable`, `clone()` makes a copy. Defense: override `clone()` to throw.

The clean shortcut for all three: use an enum singleton, which the JVM protects against reflection and serialization automatically, so you do not have to write any of these defenses.

14. How do you test components that depend on a Singleton?

Honestly, this is the pattern's biggest practical pain. A class that calls `Singleton.getInstance()` internally is hard to test because you cannot swap in a fake, the dependency is hidden and global. The way I deal with it: have the singleton implement an interface and inject it (constructor injection) rather than calling `getInstance()` inside the class, so in a test I pass a mock. So the real answer to "how do you test it" is usually "stop calling `getInstance()` directly and inject it instead", which is why people move singletons to dependency injection.

15. Would you use a Singleton for a database connection? Why or why not?

Not a raw connection, no. A single shared `Connection` is a bad idea because connections are not thread-safe and one shared connection becomes a bottleneck and a single point of failure. What you actually want is a connection pool, many connections handed out and reused, and the pool itself is often a single shared instance. So the nuance I would give: singleton the pool, never the connection. And in a Spring app I would let the framework manage that pool as a singleton bean rather than hand-rolling it.

16. Have you used a Singleton in production, which implementation, and have you replaced one with DI?

In practice the singletons I have dealt with were usually framework-managed (a Spring bean at singleton scope) rather than a hand-written `getInstance()`, because that gives you the single-instance lifecycle without the global-static downsides. When I did write one by hand, I reached for Bill Pugh (or an enum for the bulletproof case). And yes, I have replaced hand-rolled singletons with dependency injection, the reason was always testability: the moment we needed to mock the thing in a unit test, the static `getInstance()` was in the way, so injecting it behind an interface solved that. (Adapt to your own examples.)

17. If Spring already manages singleton beans, why does the Singleton pattern still exist?

Because they are not the same thing, and this catches people out. The GoF Singleton enforces one instance per JVM with a private constructor and static access, and it is self-managed global state. Spring's singleton scope means one instance per Spring container, per bean, managed by the framework and handed to you via

dependency injection, no private constructor, no static access, and you can swap it in tests. So Spring gives you the “one shared instance” benefit without the testability and coupling problems of the classic pattern. The pattern still exists for code outside a DI container, but inside one, the framework’s version is better.

18. How would you decide between a Singleton, dependency injection, and a static utility class?

The way I think about it: if the thing is pure stateless helpers (formatting, math), a static utility class is fine, no instance needed. If it has behaviour or dependencies you will want to swap or mock (a service, a repository, a client), use dependency injection, which is the default for anything with real logic. I reach for a hand-written Singleton only when I am outside a DI framework and genuinely need one shared instance. In a Spring app that middle case basically disappears, DI covers it. So: static class for stateless helpers, DI for real behaviour, classic singleton as the fallback when there is no container.

19. Is Singleton an anti-pattern, and would you keep it if redesigning Java today?

I would not call it an anti-pattern outright, but it is often misused as one. It becomes an anti-pattern when it is really a global variable in disguise: hidden mutable state that everyone reaches into, which makes code hard to test and hard to reason about. Used honestly, for something genuinely single and mostly stateless, and ideally injected behind an interface, it is fine. If I were redesigning Java today I would keep it as a documented pattern but push dependency injection as the default, because DI gives you the same “one instance” outcome without the global-state baggage. So: not evil, just overused, and usually there is a better tool. (This is a judgment question, so a defensible position with reasons is what matters.)

20. Common mistakes.

- Forgetting to make the constructor private, so callers can just new it and the guarantee is gone.
- Ignoring thread safety in lazy initialization, which races into two instances.
- Missing volatile in double-checked locking, which risks a half-constructed instance.
- Forgetting `readResolve()` when the singleton is `Serializable`, so deserialization makes a new one.
- Allowing cloning, which copies the instance (override `clone()` to throw).
- Ignoring reflection attacks on the private constructor (throw if an instance exists, or use an enum).
- Using a Singleton as a global variable for convenience, the classic misuse.
- Overusing the pattern, so shared global state spreads through the codebase.
- Using a Singleton where dependency injection fits better, which hurts testability.
- Making a mutable Singleton shared across threads without proper synchronisation.